

DiffTrace: Efficient Whole-Program Trace Analysis and Diffing

Saeed Taheri

School of Computing

University of Utah

Salt Lake City, Utah, USA

staheri@cs.utah.edu

Ian Briggs

School of Computing

University of Utah

Salt Lake City, Utah, USA

ian.briggs@gmail.com

Ganesh Gopalakrishnan

School of Computing

University of Utah

Salt Lake City, Utah, USA

ganesh@cs.utah.edu

Martin Burtscher

Department of Computer Science

Texas State University

San Marcos, Texas, USA

burtscher@cs.txstate.edu

Abstract— Abstract to be written

Index Terms— diffing, tracing, debugging

I. INTRODUCTION

When the next version of an HPC software system is created, logical errors often get introduced. To maintain productivity, designers need effective and efficient methods to locate these errors. Given the increasing use of hybrid (MPI + X) codes and library functions, errors may be introduced through a usage contract violation at any one of these interfaces. Therefore, tools that record activities at multiple APIs are necessary. Designs find most of these bugs manually, and the efficacy of a debugging tool is often measured by how well it can highlight the salient differences between the executions of two versions of software. Given the huge number of things that could be different – individual iterative patterns of function calls, groups of functions calls, or even specific instruction types (e.g., non-vectorized versus vectorized floating-point dot vector loops) – designers cannot often afford to rerun the application multiple times to collect each facet of behavior separately. These issues are well summarized in many recent studies [1]

One of the major challenges of HPC debugging is the huge diversity of the applications, which encompass domains such as computational chemistry, molecular dynamics, and climate simulation. In addition, there are many types of possible “bugs” or, more precisely, errors. An **error** may be a deadlock or a resource leak. These errors may be caused by different **faults**: an unexpected message reordering rule (for a deadlock) or a forgotten free statement (for a resource leak). There exists a collection of scenarios in which a bug can be introduced: when developing a brand new application, optimizing an existing application, upscaling an application, porting to a new platform, changing the compiler, or even changing compiler flags. Unlike traditional software, there are hardly any bug-repositories, collection of trace data or debugging-purpose benchmarks in HPC community. The heterogeneous nature of HPC bugs make developers come up with their own solutions to resolve specific class of bugs on specific architecture or platforms that are not usable on other [1].

When a failure occurs (e.g., deadlock or crash) or the application outputs an unexpected result, it is not economic to

rerun the application and consume resources to reproduce the failure. In addition, HPC bugs might not be reproducible due to non-deterministic behavior of most of HPC applications. Also the failure might be caused by a bug present at different APIs, system levels or network, thus multiple reruns might be needed to locate the buggy area. In our previous work[2], we have introduced ParLOT that collects whole program function call traces efficiently using dynamic binary instrumentation. ParLOT captures function calls and returns at different levels (e.g., source-code and library) and incrementally compress them on-the-fly, resulting in low runtime overhead and significantly low required bandwidth, preserving the whole-program dynamic behavior for offline analysis.

In the current work, we introduce DiffTrace, a tool-chain that post-mortem analyze ParLOT traces in order to supply developers with information about dynamic behavior of HPC applications at different levels towards debugging. Topology of HPC tasks on both distributed and shared memory often follow a “symmetric” control flow such as SPMD, master/worker and odd/even where multiple tasks contain *similar* events in their control flow. HPC bugs often manifest themselves as divergence in the control flow of processes comparing to what was expected. In other words, HPC bugs violates the rule of “symmetric” and “similar” control flow of one or more thread/process in typical HPC applications based on the original topology of the application. We believe that finding the dissimilarities among traces is the essential initial step towards finding the bug manifestation, and consequently the bug root cause.

Large-scale HPC application execution would result in thousands of ParLOT trace files due to execution of thousands of processes and threads. Since HPC applications spend most of their time in an outer main loop, every single trace file also may contain million-long sequence of trace entries (i.e., function calls and returns). Finding the bug manifestation (i.e., dissimilarities caused by the bug) among large number of long traces is the problem of finding the needle in the haystack.

Decompressing ParLOT traces collected from long-running large-scale HPC applications for offline analysis produce overwhelming amount of data. However, missing any piece of collected data may result in losing key information about the application behavior. We propose a variation of NLR (Nested

Loop Recognition) algorithm [3] that takes a sequence of trace entries as input and by recursively detecting repetitive patterns, re-compresses traces into “iterative sets” in a lossless fashion (intra-trace compression). Analyzing the application execution as a whole is another goal that we are pursuing in this work. By extracting *attributes* from pre-processed traces, we inject them a concept hierarchy data structure called Concept Lattice [4]. Concept lattices give us the capability of reducing the search space from thousands of instances to just a few *equivalent behavior classes of traces* by measuring the similarity of traces[5]

** TODO: Highlights of results obtained as a result of the above thinking should be here. THis typically comes before ROADMAP of paper.

In summary, here are our main contributions:

- we have a powerful combination of ideas to locate bugs
- A variation of NLR algorithm to compress traces in lossless fashion for easier analysis and detecting (broken) loop structures
- FCA-based clustering of similar behavior traces, efficient,
- Ranking system based on delta-sim
- Visualization diffNLR

The rest of the paper is as follows:

- Sec 2: Background
- Sec 3: Related Work
- Sec 4: DiffTrace Components and Design
- Sec 5: DiffTrace Evaluation and Experiments
- Sec 6: Discussion, Limitations, Conclusion, Future Work

II. BACKGROUND

There are two major phases in any “Program Understanding” tool: *data collection* and *data analysis*. To understand the runtime behavior of applications, an efficient tracing mechanism is required to collect informative data during execution of the application. Upon failure or observing unexpected behavior of the program (e.g., wrong answer), studying collected execution data would reveal insight about how program dynamically behaved and what went wrong. In this section, we explain our methodology of data collection and data analysis towards debugging and locating potential root causes of unexpected behavior. The main source of whole-program dynamic behavior is provided by ParLOT, a dynamic binary tracing tool that captures all function calls and returns and compress them on the fly (section II-A). After pre-processing ParLOT traces, a loop-detection-based lossless data reduction mechanism applies to each trace to simplify collected data and reflect facts about loop structures (section II-B). Whole-program analysis in HPC applications only makes sense when the analysis includes cross-thread and cross-process as well as analysis of sequential control flow of every single running thread. Inspired by many works[6] [5] [7][8], we have used FCA[4] to integrate collected data into a single data structure as a whole, and extract valuable information about different aspect of the execution (section II-C) The major advantage of FCA, is that we can extract a full pair-wise similarity

score matrix for all traces of a single execution in an efficient and scalable way, based on attributes that we extract from pre-processed traces. Relying on similarities of traces, we classify traces (i.e., threads) into equivalent-behavior groups. This way we reduce the search space from thousands of long traces to just a few classes of simplified trace representation. Any abnormal behavior of a thread (e.g., outlier) or group of threads can be detected (section II-E), and for more-in-depth analysis, any pair of traces can be studied with respect to their entry orders and differences using a visualization (section II-D).

Section II-F discusses some additional remarks (limitations, effectiveness of our approach)

A. ParLOT: Efficient Trace Collection

The final executable of real HPC applications are often a production of a large code base and a complex build system with numerous dependencies and libraries. Injecting instrumentation code to the source code, as in traditional tools like [????], is not feasible in HPC space. Also recompilation of the application with tools’ compile-wrappers, as in TAU[9] and Score-p[10], may break the build system. Also instrumentation and tracing mechanism of existing tools are often dependent to other libraries that are need to be present on the supercomputer for trace collection. [[Example: STAT[11] and AutomaDeD[12] that requires Dyninst[13] for instrumentation and MRNet[14] and TBON [15]]] To overcome the trade-off of comprehensive data collection while adding low time and space overhead, HPC program analysis tools often sacrifice one for the other. However, ParLOT collects whole-program function call traces at as low as library level, while incrementally compressing traces on-the-fly and leave majority of the system bandwidth for the application. ParLOT collects *whole* program function call traces with the mindset of *paying a little upfront and save resource and time cost of reproducing the bug*. ParLOT instruments the entry and exit point of each function in the binary using Pin[16] (fig. 1). Each ParLOT Trace contain full sequence of function calls and returns for every single thread that running the application code, reflecting the dynamic control flow and call-stack of the application individually. Here we define ParLOT Trace, as we refer to it in the paper:

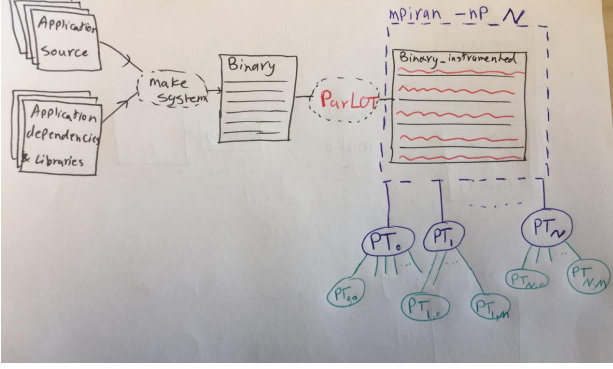
Definition 1: ParLOT Trace (PT) A ParLOT Trace (PT) is a sequence of ordered integers $\langle f_i, \dots, f_j \rangle$ where f_0 is *return* and f_k is the id of *function* k ($k! = 0$).

Note that the PT is the pre-processed (decompressed and filtered) version of immediate ParLOT traces. The fresh output of ParLOT traces are highly compressed byte-codes. Also Note that $PT_{p,t}$ refers to the PT that belongs to process p and thread t of that process.

B. Loop structure detection

HPC applications and resources are in interest of scientists and engineers for simulating *iterative* kernels. Computer simulation of fluid dynamics, partial differential equations, Gauss-

Figure 1. ParLOT Overview



Seidel method and finite element methods in form of stencil codes do include a main outer loop that iterates over some elements (i.e., timesteps) and updates elements. This character of typical HPC applications make PT lengths too long (millions) with much smaller number (hundreds) of distinct elements (i.e., function IDs). We propose a representation of PT elements (intra-PT compression) in form of *loop structures*, such that PT = sequence of repetitive patterns (i.e, loops). In other words, each PT is a sequence of *Loop Bodies (LB)* that repeated *Loop Count (LC)* times, consecutively.

1) *Loops definition*: According to Makoto Kobayashi[17] definition of loops, an occurrence of a *loop* is defined as a *sequence of elements* in which a particular sequence of *distinct elements* (called the *cycle* of the loop) is successively repeated. Later Alain Ketterlin in [3] expanded this definition to numerical values for compressing and predicting memory access addresses and designed Nested Loop Recognition (NLR) algorithm. The basic idea behind NLR algorithm is that a linear function model can be extracted from the linear progression in a sequence of numbers and these linear functions form a tree in which depth of each node is the depth of *nested* loop (e.g., most outer loop's function is the root of the tree with depth 0). We have modified NLR algorithm to make it suitable for PTs. Each repetitive pattern and its frequency of consecutive appearances would be compressed to a single *Loop Structure (LS)* entry.

Definition 2: Loop Structure $LS = LB \wedge LC$ where $LB = \langle pt_i, \dots, pt_j \rangle$ ($0 \leq i < j < len(PT)$) that occur LC times is in a sub-sequence $\langle pt_i, \dots, pt_k \rangle$ ($k = 3j, k < len(PT)$). By converting each PT to a sequence of LS_i , we reduce the length of PT by a factor of $\sum_i len(LB_i) * LC_i$. Later we will explain how this lossless representation of PTs eases the process of diffing between a pair of PTs.

C. Equivalencing behavior via FCA

To Reduce the search space from thousands of PTs to just a few groups of equivalent PTs (i.e., inter-PT compression) only requires a similarity measure based on a call matrix, but also requires a scheme that is efficient even for large process counts. Since a pair-wise comparison of all processes is highly inefficient, we use *concept lattices* that stem from *formal concept analysis (FCA)*[4] to store and compute groups of

Table I
CONTEXT

	MPI_Init()	MPI_Comm_Size()	MPI_Comm_Rank()	MPI_Send()	MPI_Recv()	MPI_Finalize()
Rank 0	x	x	x		x	x
Rank 1	x	x	x	x		x
Rank 2	x	x	x	x		x
Rank 3	x	x	x	x		x

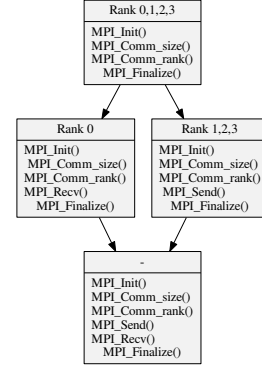


Figure 2. Sample Concept Lattice from Obj-Atr Context in table II-C

similar PTs. A concept lattice is based on a *formal context*[4], which is a triple (O, A, I) , where O is a set of **objects**, A a set of **attributes**, and $I \subseteq O \times A$ an incidence relation. The incidence relation associates each object with a set of attributes. Due to its valuable properties, especially its *partial order*, FCA has been used widely in computer science fields from machine learning and data mining [7] to distributed systems [8]. However, since we are only interested in grouping similar PTs in this work, we only take advantage of similarity measures [5] of concept lattices, and left rest of its properties for our future work. Due to typical HPC application common topologies such as SPMD, master/worker and odd/even where multiple processes/threads behave similarly, our experiments show that large number of PTs can be reduced to just a few groups.

```
main() {
    int rank;
    int src;
    MPI_Init()
    MPI_Comm_size(MPI_COMM_WORLD)
    MPI_Comm_rank(MPI_COMM_WORLD, &rank)
    if (rank != 0) {
        MPI_Send(0) // Send to rank 0
    } else { /* rank = 0
        MPI_Recv(1) // Receive from rank 1
        MPI_Recv(2) // Receive from rank 2
        MPI_Recv(3) // Receive from rank 3
    }
    MPI_Finalize()
}
```




Figure 3. Concept Lattice with reduced labels




Figure 4. Pair-wise Jaccard Similarity Matrix (JSM) of MPI processes in Sample code

An example with

- Source code [done]
- Simple attributes and Context [done]
- Concept Lattice [done]
- Jaccard Similarity (heatmap or matrix) [done]
- How we obtained JSM elements
- Visualization of LCA [not yet]
- CL construction [18] [19]
- Attribute Table

D. *diffNLR*: Reflecting differences

-From fig we can see that PT0 and PT1 are in different groups with similarity of XX.

Inspired by *diff* original algorithm[20] that has bin used in Git and GNU Diff, we visualize the differences of a pair of PT as shown in fig 5.

This visualization reflects of the differences of **occurrences** of PT elements and their **orders**.

In section V we show how this visualization can help us locating the points of divergence in PTs, and potential bug manifestation and root cause.




Figure 5. Sample *diffNLR*

E. Experimental Methodology

- Candidates to look at and observe their *diffNLR*: $(PT_i, PT_j) = \text{Max}|JSM[i][j](buggy) - JSM[i][j](bugfree)|$

- It means the similarity of (PT_i, PT_j) of buggy execution have the highest difference with its corresponding similarity of the bug-free execution.

- Figures are showing two *diffNLR* side by side, one is $diffNLR(i, j)_{buggy}$ and the other one is $diffNLR(i, j)_{bugfree}$ where i, j obtained from candidate tables.

- Two additional tables that show the LB (loop body) of each detected loop. L_i

F. Additional Remarks

III. RELATED WORK

A. Program Understanding

- Score-P [10]
- TAU [9]
- ScalaTrace: Scalable compression and replay of communication traces for HPC [21]
- Barrier Matching for Programs with Textually unaligned barriers [22]
- Pivot Tracing: Dynamic causal monitoring for distributed systems - Johnathan mace [23]
- Automated Charecterization of parallel application communication patterns [24]
- Problem Diagnosis in Large Scale Computing environments [25]
- Probablistic diagnosis of performance faults in large-scale parallel applications [26]
- detecting patterns in MPI communication traces - robert preissl [27]

Include a table with all filters and their regular expressions

1) General Filters:

- Returns
- .plt
- Memory
- Network
- Polling
- String
- Customize
- IncludeEverything

2) Target Filters:

- MPI_
- MPIall
- MPI_ Collectives
- MPI_ Send/Recv
- OMPall
- OMPcritical
- OMPmutex

E. Nested Loop Recognition

1) Background:

2) Implementation:

F. Concept Lattice Analysis

1) Background:

- FCA (formal concept analysis) background and citations
- FCA applications in all areas
- FCA applications in Data Mining and Information Retrieval
- FCA applications in distributed systems (Garg's work)
- intro to Concept, Object, Attribute and other definitions

2) Objects/Attributes: Mapping of Object/Attribute (general) to Trace/Attribute (clTrace)

What do we expect to gain by doing so

- Single entity represents the whole execution of HPC application (can be used as signature/model in ML)
- Classifying similar behavior objects(traces)
- Efficient Incremental CL building makes it scalable
- Efficient full pair-wise Jaccard Similarity Matrix extraction

3) CL generation:

- background
- current approach

4) Jaccard Similarity Matrix:

- background
- LCA
- Benefits

G. diffNLR

- motivation
- diff algorithm
- visualization

H. FP-Trace

V. EXPERIMENTAL STUDIES

A. Case Study: ILCS-TSP

Here is the ILCS framework pseudo-code. User needs to write CPU_Init(), CPU_Exec() and CPU_Output().

```
1  int main(argc, argv){
2  MPI_Init();
3  MPI_Comm_size()
4  MPI_Comm_rank(my_rank)
5  //Figuring local number of CPUs
6  MPI_Reduce() // Figuring global number of CPUs
7  CPU_Init();
8  //For storing local champion results
9  champ[CPUs] = malloc();
10 MPI_Barrier();
11 #pragma omp parallel num_threads(CPUs+1)
12 {
13     rank = omp_get_thread_num()
14     if (rank == 0){ //communication thread
15         do{
16             //Find and report the thread with
17             //local champion, global champion
18             MPI_AllReduce();
19             //Find and report the process with
20             //global champion
21             MPI_AllReduce();
22             //The process with the global champion
23             //copy its results to bcast_buffer
24             if (my_rank == global_champion){
25                 #pragma omp cirtical
26                 memcpy(bcast_buffer, local_champ)
27             }
28             //Broadcast the champion
29             MPI_Bcast(bcast_buffer)
30         } while (no_change_threshold);
31         cont=0 // signal worker threads to stop
32     } else{ // worker threads
33         while(cont){
34             //Calculate Seed
35             local_result = CPU_Exec()
36             if (local_result < champ[rank]){
37                 #pragma omp cirtical
38                 memcpy(champ[rank], local_result)
39             }
40         }
41     }
42 }
43 //Find and report the thread with
44 //local champion, global champion
45 MPI_AllReduce();
46 //Find and report the process with
47 //global champion
48 MPI_AllReduce();
```

```

49 // The process with the global champion
50 // copy its results to bcast_buffer
51 if (my_rank == global_champion){
52     #pragma omp cirtical
53     memcpy( bcast_buffer , local_champ )
54 }
55 //Broadcast the champion
56 MPI_Bcast( bcast_buffer )
57 if ( my_rank==0){
58     CPU_Output( champ )
59 }
60 MPI_Finalize()
61 }
62
63 /* User code for TSP problem */
64
65 CPU_Init(){
66     // Read In data from cities
67     // Calculate distances
68     // Return data structure to store champion
69 }
70
71 CPU_Exec(){
72     // Find local champions (TSP tours)
73 }
74
75 CPU_Output(){
76     // Output champion
77 }

```


Table ?? describes the bug that I injected to ILCS-TSP

1) 3: Missing Critical Section one thread in on process: I planted the bug (missing critical section) in process 2

VI. DISCUSSION AND FUTURE WORK

REFERENCES

- [1] G. Gopalakrishnan, P. D. Hovland, C. Iancu, S. Krishnamoorthy, I. Laguna, R. A. Lethin, K. Sen, S. F. Siegel, and A. Solar-Lezama, "Report of the HPC correctness summit, jan 25-26, 2017, washington, DC," *CoRR*, vol. abs/1705.07478, 2017. [Online]. Available: <http://arxiv.org/abs/1705.07478>
- [2] S. Taheri, S. Devale, G. Gopalakrishnan, and M. Burtcher, "Parlot: Efficient whole-program call tracing for hpc applications," in *7th Workshop on Extreme-Scale Programming Tools, ESPT@SC 2018, Dallas, TX, USA, November 16, 2018*.
- [3] A. Ketterlin and P. Clauss, "Prediction and trace compression of data access addresses through nested loop recognition," in *Proceedings of the 6th Annual IEEE/ACM International Symposium on Code Generation and Optimization*, ser. CGO '08. New York, NY, USA: ACM, 2008, pp. 94–103. [Online]. Available: <http://doi.acm.org/10.1145/1356058.1356071>
- [4] B. Ganter and R. Wille, *Formal Concept Analysis: Mathematical Foundations*, 1st ed. Secaucus, NJ, USA: Springer-Verlag New York, Inc., 1997.
- [5] F. Alqadah and R. Bhatnagar, "Similarity measures in formal concept analysis," *Annals of Mathematics and Artificial Intelligence*, vol. 61, no. 3, pp. 245–256, Mar 2011. [Online]. Available: <https://doi.org/10.1007/s10472-011-9257-7>
- [6] M. Weber, R. Brendel, T. Hilbrich, K. Mohror, M. Schulz, and H. Brunst, "Structural Clustering: A New Approach to Support Performance Analysis at Scale." *IEEE*, May 2016, pp. 484–493. [Online]. Available: <http://ieeexplore.ieee.org/document/7516045/>
- [7] D. I. Ignatov, "Introduction to formal concept analysis and its applications in information retrieval and related fields," *CoRR*, vol. abs/1703.02819, 2017. [Online]. Available: <http://arxiv.org/abs/1703.02819>
- [8] V. K. Garg, "Maximal antichain lattice algorithms for distributed computations," in *Distributed Computing and Networking*, D. Frey, M. Raynal, S. Sarkar, R. K. Shyamasundar, and P. Sinha, Eds. Berlin, Heidelberg: Springer Berlin Heidelberg, 2013, pp. 240–254.
- [9] S. S. Shende and A. D. Malony, "The Tau parallel performance system," *International Journal on High Performance Computer Applications*, vol. 20, pp. 287–311, May 2006. [Online]. Available: <http://portal.acm.org/citation.cfm?id=1125980.1125982>
- [10] A. Knüpfer, C. Rössel, D. an Mey, S. Biersdorff, K. Diethelm, D. Eschweiler, M. Geimer, M. Gerndt, D. Lorenz, A. D. Malony, W. E. Nagel, Y. Oleynik, P. Philippen, P. Saviankou, D. Schmidl, S. Shende, R. Tschüter, M. Wagner, B. Wesarg, and F. Wolf, "Score-p: A joint performance measurement run-time infrastructure for periscope, scalasca, tau, and vampir," in *Tools for High Performance Computing 2011 - Proceedings of the 5th International Workshop on Parallel Tools for High Performance Computing, ZIH, Dresden, September 2011*, 2011, pp. 79–91.
- [11] D. H. Ahn, B. R. de Supinski, I. Laguna, G. L. Lee, B. Liblit, B. P. Miller, and M. Schulz, "Scalable temporal order analysis for large scale debugging," in *Proceedings of the Conference on High Performance Computing Networking, Storage and Analysis*, Nov 2009, pp. 1–11.
- [12] I. Laguna, T. Gambin, B. R. de Supinski, S. Bagchi, G. Bronevetsky, D. H. Anh, M. Schulz, and B. Rountree, "Large scale debugging of parallel tasks with automated," in *Proceedings of 2011 International Conference for High Performance Computing, Networking, Storage and Analysis*, ser. SC '11. New York, NY, USA: ACM, 2011, pp. 50:1–50:10. [Online]. Available: <http://doi.acm.org/10.1145/2063384.2063451>
- [13] B. P. Miller, M. D. Callaghan, J. M. Cargille, J. K. Hollingsworth, R. B. Irvin, K. L. Karavanic, K. Kunchithapadam, and T. Newhall, "The paradigm parallel performance measurement tool," *IEEE Computer*, vol. 28, no. 11, pp. 37–46, 1995. [Online]. Available: <https://doi.org/10.1109/2.471178>
- [14] P. C. Roth, D. C. Arnold, and B. P. Miller, "Mrnet: A software-based multicast/reduction network for scalable tools," in *Supercomputing, 2003 ACM/IEEE Conference*, Nov 2003, pp. 21–21.
- [15] T. Hilbrich, B. R. de Supinski, F. Hänsel, M. S. Müller, M. Schulz, and W. E. Nagel, "Runtime mpi collective checking with tree-based overlay networks," in *Proceedings of the 20th European MPI Users' Group Meeting*, ser. EuroMPI '13. New York, NY, USA: ACM, 2013, pp. 129–134. [Online]. Available: <http://doi.acm.org/10.1145/2488551.2488570>
- [16] C.-K. Luk, R. Cohn, R. Muth, H. Patil, A. Klauser, G. Lowney, S. Wallace, V. J. Reddi, and K. Hazelwood, "Pin: Building customized program analysis tools with dynamic instrumentation," in *Proceedings of the 2005 ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '05. New York, NY, USA: ACM, 2005, pp. 190–200.
- [17] "Dynamic characteristics of loops," *IEEE Transactions on Computers*, vol. C-33, no. 2, pp. 125–132, Feb 1984.
- [18] R. Godin, R. Missaoui, and H. Alaoui, "Incremental concept formation algorithms based on galois (concept) lattices," *Computational Intelligence*, vol. 11, no. 2, pp. 246–267.
- [19] M. A. Bender, M. Farach-Colton, G. Pemmasani, S. Skiena, and P. Sumazin, "Lowest common ancestors in trees and directed acyclic graphs," *Journal of Algorithms*, vol. 57, no. 2, pp. 75 – 94, 2005. [Online]. Available: <http://www.sciencedirect.com/science/article/pii/S0196677405000854>
- [20] E. W. Myers, "An O(ND) difference algorithm and its variations," *Algorithmica*, vol. 1, no. 2, pp. 251–266, 1986. [Online]. Available: <https://doi.org/10.1007/BF01840446>
- [21] M. Noeth, P. Ratn, F. Mueller, M. Schulz, and B. R. de Supinski, "Scalatrace: Scalable compression and replay of communication traces for high-performance computing," *Journal of Parallel and Distributed Computing*, vol. 69, no. 8, pp. 696 – 710, 2009, best Paper Awards: 21st International Parallel and Distributed Processing Symposium (IPDPS 2007).
- [22] Y. Zhang and E. Duesterwald, "Barrier matching for programs with textually unaligned barriers," in *Proceedings of the 12th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, ser. PPoPP '07. New York, NY, USA: ACM, 2007, pp. 194–204. [Online]. Available: <http://doi.acm.org/10.1145/1229428.1229472>
- [23] J. Mace, R. Roelke, and R. Fonseca, "Pivot tracing: Dynamic causal monitoring for distributed systems," *ACM Trans. Comput. Syst.*, vol. 35, no. 4, pp. 11:1–11:28, Dec. 2018. [Online]. Available: <http://doi.acm.org/10.1145/3208104>
- [24] P. C. Roth, J. S. Meredith, and J. S. Vetter, "Automated characterization of parallel application communication patterns," in *Proceedings of the 24th International Symposium on High-Performance Parallel and Distributed Computing*, ser. HPDC '15. New York, NY, USA: ACM, 2015, pp. 73–84. [Online]. Available: <http://doi.acm.org/10.1145/2749246.2749278>
- [25] A. V. Mirgorodskiy, N. Maruyama, and B. P. Miller, "Problem diagnosis in large-scale computing environments," in *SC '06: Proceedings of the 2006 ACM/IEEE Conference on Supercomputing*, Nov 2006, pp. 11–11.
- [26] I. Laguna, D. H. Ahn, B. R. de Supinski, S. Bagchi, and T. Gambin, "Probabilistic diagnosis of performance faults in large-scale parallel applications," in *Proceedings of the 21st International Conference on Parallel Architectures and Compilation Techniques*, ser. PACT '12. New York, NY, USA: ACM, 2012, pp. 213–222. [Online]. Available: <http://doi.acm.org/10.1145/2370816.2370848>
- [27] R. Preissl, T. Köckerbauer, M. Schulz, D. Kranzlmüller, B. R. de Supinski, and D. J. Quinlan, "Detecting patterns in mpi communication traces," *2008 37th International Conference on Parallel Processing*, pp. 230–237, 2008.
- [28] B. Liu and J. Huang, "D4: Fast concurrency debugging with parallel differential analysis," in *Proceedings of the 39th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI 2018. New York, NY, USA: ACM, 2018, pp. 359–373. [Online]. Available: <http://doi.acm.org/10.1145/3192366.3192390>
- [29] B. Krammer, M. MÄßler, and M. Resch, "Mpi application development using the analysis tool marmot," vol. 3038, 12 2004, pp. 464–471.
- [30] A. Droste, M. Kuhn, and T. Ludwig, "Mpi-checker: Static analysis for mpi," in *Proceedings of the Second Workshop on the LLVM Compiler Infrastructure in HPC*, ser. LLVM '15. New York, NY, USA: ACM, 2015, pp. 3:1–3:10. [Online]. Available: <http://doi.acm.org/10.1145/2833157.2833159>
- [31] Q. Gao, F. Qin, and D. K. Panda, "Dmtracker: finding bugs in large-scale parallel programs by detecting anomaly in data movements," in *SC '07: Proceedings of the 2007 ACM/IEEE Conference on Supercomputing*, Nov 2007, pp. 1–12.

Table II
INJECTED BUGS TO ILCS-TSP

ID	Level	Bugs	Description
1	MPI	allRed1wrgOp-1-all-x	Different operation (MPI_MAX) in only one process (buggyProc = 2) for MPI_ALLREDUCE() in Line 21
2		allRed1wrgSize-1-all-x	Wrong size in only one process (buggyProc = 2) for MPI_ALLREDUCE() in Line 21
3		allRed1wrgSize-all-all-x	Wrong Size in all processes for MPI_ALLREDUCE() in Line 21
4		allRed2wrgOp-1-all-x	Different operation (MPI_MAX) in only one (buggyProc) for first MPI_ALLREDUCE() – L277:ilcsTSP.c
5		allRed2wrgSize-1-all-x	Wrong size in only one (buggyProc) for first MPI_ALLREDUCE() – L277:ilcsTSP.c
6		allRed2wrgSize-all-all-x	Wrong Size in all processes for second MPI_ALLREDUCE() – L277:ilcsTSP.c
7		bcastWrgSize-1-all-x	Wrong Size in only one (buggyProc) of MPI_Bcast() – L290:ilcsTSP.c
8		bcastWrgSize-all-all-x	Wrong Size n all processes for MPI_Bcast() – L240:ilcsTSP.c
9	OMP	misCrit-1-1-x	Missing Critical Section in buggyProc and buggyThread – L170:ilcsTSP.c
10		misCrit-all-1-x	Missing Critical Section in buggyThread and all procoesses – L170:ilcsTSP.c
11		misCrit-1-all-x	Missing Critical Section in buggyProc and all threads – L170:ilcsTSP.c
12		misCrit-all-all-x	Missing Critical Section in all procs and threads – L170:ilcsTSP.c
13		misCrit2-1-1-x	Missing Critical Section in buggyProc and buggyThread – L230:ilcsTSP.c
14		misCrit2-all-1-x	Missing Critical Section in buggyThread – L230:ilcsTSP.c
15		misCrit2-1-all-x	Missing Critical Section in buggyProc and all threads – L230:ilcsTSP.c
16		misCrit2-all-all-x	Missing Critical Section in all procs and threads – L230:ilcsTSP.c
17		misCrit3-1-all-x	Missing Critical Section in buggyProc and all threads – L280:ilcsTSP.c
18		misCrit3-all-all-x	Missing Critical Section in all procs and threads – L280:ilcsTSP.c
19	General	infLoop-1-1-1	Injected an infinite loop after CPU_EXEC() in buggyProc,buggyThread & buggyIter L164:ilcsTSP.c

45	(7)11.mem.ompcrit.cust.0K10	sing.actual	1: (3_0,4_0):0.75 1: (2_1,3_1):0.00 1: (3_2,6_2):0.57 1: (1_3,2_3):0.89
			2: (2_0,3_0):0.75 2: (0_1,4_1):0.00 2: (3_2,5_2):0.57 2: (0_3,7_3):0.89
			3: (1_0,6_0):0.75 3: (0_1,5_1):0.00 3: (2_2,4_2):0.50 3: (4_3,6_3):0.50
46	(7)11.mem.ompcrit.cust.0K10	sing.log10	1: (3_0,4_0):0.75 1: (2_1,3_1):0.00 1: (0_2,2_2):0.33 1: (1_3,2_3):0.33
			2: (2_0,3_0):0.75 2: (0_1,4_1):0.00 2: (3_2,7_2):0.33 2: (6_3,7_3):0.33
			3: (1_0,6_0):0.75 3: (0_1,5_1):0.00 3: (3_2,6_2):0.33 3: (1_3,4_3):0.20
47	(7)11.mem.ompcrit.cust.0K10	sing.orig	1: (3_0,4_0):0.75 1: (2_1,3_1):0.00 1: (0_2,2_2):0.33 1: (1_3,2_3):0.33
			2: (2_0,3_0):0.75 2: (0_1,4_1):0.00 2: (3_2,7_2):0.33 2: (6_3,7_3):0.33
			3: (1_0,6_0):0.75 3: (0_1,5_1):0.00 3: (3_2,6_2):0.33 3: (1_3,4_3):0.20

Figure 7. Part of ranking table for MisCrit 1-1

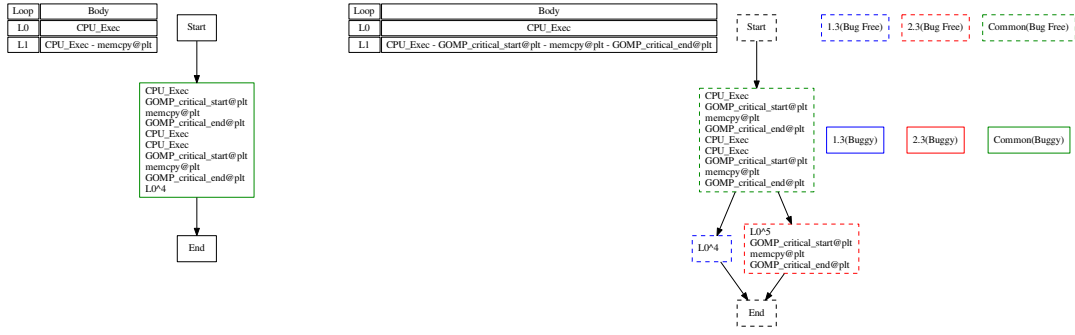


Figure 8. diffNLR of process 1 thread 3 and process 2 thread 3 buggy(missing critical section) vs. bug-free

- [32] Z. Chen, X. Li, J. Chen, H. Zhong, and F. Qin, “Syncchecker: Detecting synchronization errors between mpi applications and libraries,” in *2012 IEEE 26th International Parallel and Distributed Processing Symposium*, May 2012, pp. 342–353.
- [33] N. Maruyama and S. Matsuoka, “Model-based fault localization in large-scale computing systems,” in *2008 IEEE International Symposium on Parallel and Distributed Processing*, April 2008, pp. 1–12.
- [34] I. Beschastnikh, J. Abrahamson, Y. Brun, and M. D. Ernst, “Synoptic: Studying logged behavior with inferred models,” in *Proceedings of the 19th ACM SIGSOFT Symposium and the 13th European Conference on Foundations of Software Engineering*, ser. ESEC/FSE ’11. New York, NY, USA: ACM, 2011, pp. 448–451. [Online]. Available: <http://doi.acm.org/10.1145/2025113.2025188>
- [35] I. Beschastnikh, Y. Brun, M. D. Ernst, A. Krishnamurthy, and T. E. Anderson, “Mining temporal invariants from partially ordered logs,” in *Managing Large-scale Systems via the Analysis of System Logs and the Application of Machine Learning Techniques*, ser. SLAML ’11. New York, NY, USA: ACM, 2011, pp. 3:1–3:10. [Online]. Available: <http://doi.acm.org/10.1145/2038633.2038636>
- [36] D. H. Ahn, B. R. de Supinski, I. Laguna, G. L. Lee, B. Liblit, B. P. Miller, and M. Schulz, “Scalable temporal order analysis for large scale debugging,” in *Proceedings of the Conference on High Performance Computing Networking, Storage and Analysis*, ser. SC ’09. New York, NY, USA: ACM, 2009, pp. 44:1–44:11. [Online]. Available: <http://doi.acm.org/10.1145/1654059.1654104>
- [37] S. Grant, H. Cech, and I. Beschastnikh, “Inferring and asserting distributed system invariants,” in *Proceedings of the 40th International Conference on Software Engineering*, ser. ICSE ’18. New York, NY, USA: ACM, 2018, pp. 1149–1159. [Online]. Available: <http://doi.acm.org/10.1145/3180155.3180199>
- [38] S. Mitra, I. Laguna, D. H. Ahn, S. Bagchi, M. Schulz, and T. Gamblin, “Accurate application progress analysis for large-scale parallel debugging,” in *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI ’14. New York, NY, USA: ACM, 2014, pp. 193–203. [Online]. Available: <http://doi.acm.org/10.1145/2594291.2594336>
- [39] G. Bronevetsky, I. Laguna, S. Bagchi, B. R. de Supinski, D. H. Ahn, and M. Schulz, “Automated: Automata-based debugging for dissimilar par-

- allel tasks,” in *2010 IEEE/IFIP International Conference on Dependable Systems Networks (DSN)*, June 2010, pp. 231–240.
- [40] I. Beschastnikh, Y. Brun, M. D. Ernst, and A. Krishnamurthy, “Inferring models of concurrent systems from logs of their behavior with csight,” in *Proceedings of the 36th International Conference on Software Engineering*, ser. ICSE 2014. New York, NY, USA: ACM, 2014, pp. 468–479. [Online]. Available: <http://doi.acm.org/10.1145/2568225.2568246>
 - [41] M. Weber, R. Brendel, and H. Brunst, “Trace file comparison with a hierarchical sequence alignment algorithm,” in *2012 IEEE 10th International Symposium on Parallel and Distributed Processing with Applications*, July 2012, pp. 247–254.
 - [42] B. Liblit and A. Aiken, “Building a better backtrace: Techniques for postmortem program analysis,” Berkeley, CA, USA, Tech. Rep., 2002.
 - [43] T. Sherwood, E. Perelman, G. Hamerly, and B. Calder, “Automatically characterizing large scale program behavior,” in *Proceedings of the 10th International Conference on Architectural Support for Programming Languages and Operating Systems*, ser. ASPLOS X. New York, NY, USA: ACM, 2002, pp. 45–57. [Online]. Available: <http://doi.acm.org/10.1145/605397.605403>
 - [44] K. E. Isaacs, P.-T. Bremer, I. Jusufi, T. Gamblin, A. Bhatele, M. Schulz, and B. Hamann, “Combing the communication hairball: Visualizing parallel execution traces using logical time,” *IEEE Transactions on Visualization and Computer Graphics*, vol. 20, pp. 2349–2358, 2014.
 - [45] K. E. Isaacs, A. Bhatele, J. Lifflander, D. Böhme, T. Gamblin, M. Schulz, B. Hamann, and P. Bremer, “Recovering logical structure from charm++ event traces,” in *SC ’15: Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, Nov 2015, pp. 1–12.
 - [46] I. Beschastnikh, P. Wang, Y. Brun, and M. D. Ernst, “Debugging distributed systems,” *Commun. ACM*, vol. 59, no. 8, pp. 32–37, Jul. 2016. [Online]. Available: <http://doi.acm.org/10.1145/2909480>
 - [47] M. Crochemore and W. Rytter, “Usefulness of the karp-miller-rosenberg algorithm in parallel computations on strings and arrays,” *Theoretical Computer Science*, vol. 88, no. 1, pp. 59 – 82, 1991. [Online]. Available: <http://www.sciencedirect.com/science/article/pii/030439759190073B>
 - [48] R. M. Karp, R. E. Miller, and A. L. Rosenberg, “Rapid identification of repeated patterns in strings, trees and arrays,” in *Proceedings of the Fourth Annual ACM Symposium on Theory of Computing*, ser. STOC ’72. New York, NY, USA: ACM, 1972, pp. 125–136. [Online]. Available: <http://doi.acm.org/10.1145/800152.804905>
 - [49] A. Nakamura, T. Saito, I. Takigawa, M. Kudo, and H. Mamitsuka, “Fast algorithms for finding a minimum repetition representation of strings and trees,” *Discrete Applied Mathematics*, vol. 161, no. 10, pp. 1556 – 1575, 2013. [Online]. Available: <http://www.sciencedirect.com/science/article/pii/S0166218X13000024>
 - [50] M. Crochemore and W. Rytter, *Jewels of Stringology*. World Scientific, 2002. [Online]. Available: <https://books.google.com/books?id=ipuPQgAACAAJ>
 - [51] —, *Text Algorithms*. New York, NY, USA: Oxford University Press, Inc., 1994.

APPENDIX

A. Case Study: Hybrid (MPI+OMP) Matmul

Below is the JSM of bug free version of hybrid matmul applying different filters to understand its behavior.

B. Bug: AllReduce wrong op - proc 2

Bug injected within to MPI_AllReduce() in line 21 where the local champions of all processes are about computed and their MIN is going to be distributed to all other processes. However, with injecting the wrong operation (MAX instead of MIN) to one (or more) of the processes, the logic of if statement in line 24 would be violated and the control flow of the code would not take this branch which is the the core part of ILCS (deciding champion to broadcast to other processes).

After collecting traces from bug-free version and this buggy version of ILCS-TSP, and after applying filters and detecting loops, my ranking system, with the highest possible score, recommends to look at figure 13 diffNLR. This diffNLR shows that in the bug free version, we are expecting that at least one of the processes (in this case process 6 to win the championship, copy the result of that process to broadcast_buffer and broadcast it to all other processes. However, in the buggy version, the champion never gets updated (never goes through OMP critical section).

The recommendation system that I designed tries to find the pairs of traces that their similarities changed the most. In other word, this recommendation system wants to suggest top diffNLRs that are candidates to show the impact of the bug based on trace similarities. (table

C. Bug: AllReduce wrong size - one and all

This bug made the code crash with showing an error in MPI_Bcast(), although the bug was injected to MPI_AllReduce(). Figure 14 $CMPdiffNLR(P_{2.0}, P_{6.0})$ and 18 $CMPdiffNLR(P_{1.0}, P_{6.0})$, where $CMPdiffNLR(P_i, P_j)$ shows the comparison of

$$diffNLR(P_{i.buggy}, P_{j.buggy})$$

and

$$diffNLR(P_{i.bugFree}, P_{j.bugFree})$$

Ranking table is showing all 1.00 since most of the traces did not go through.

D. Bug: Missing Critical Section one thread in on process



Figure 9. JSM of bug-free version of hybrid Matmul (all functions)



Figure 10. JSM of bug-free version of hybrid Matmul (MPI and OMP functions)




Figure 11. JSM of bug-free version of hybrid Matmul (MPI Only)

	Filter	Attributes	Top Thread 0	Top Thread 1	Top Thread 2	Top Thread 3	Top Thread 4
0	(14)11.nen.rpicol.ompcrit.cust.0K30	doub.actual	1:(1,0,0) 1.00 2:(2,0,0) 1.00 3:(5,0,0) 1.00	1:(2,1,3,1):0.00 2:(0,1,4,1):0.00 3:(0,1,5,1):0.00	1:(1,2,6,2):0.57 2:(3,2,6,2):0.42 3:(1,2,3,2):0.34	1:(1,3,4,3):0.47 2:(6,3,7,3):0.30 3:(0,3,2,3):0.27	1:(6,4,7,4):0.71 2:(4,4,6,4):0.69 3:(0,4,6,4):0.69
1	(14)11.nen.rpicol.ompcrit.cust.0K30	doub.log10	1:(1,0,0) 1.00 2:(2,0,0) 1.00 3:(5,0,0) 1.00	1:(2,1,3,1):0.00 2:(0,1,4,1):0.00 3:(0,1,5,1):0.00	1:(3,2,6,2):0.36 2:(0,2,2,2):0.38 3:(1,2,3,2):0.29	1:(1,3,6,3):0.62 2:(0,3,5,3):0.29 3:(1,3,4,3):0.26	1:(1,4,5,4):0.62 2:(0,4,5,4):0.39 3:(0,4,4,4):0.33
2	(14)11.nen.rpicol.ompcrit.cust.0K30	doub.orig	1:(1,0,0) 1.00 2:(2,0,0) 1.00 3:(5,0,0) 1.00	1:(2,1,3,1):0.00 2:(0,1,4,1):0.00 3:(0,1,5,1):0.00	1:(3,2,6,2):0.36 2:(0,2,2,2):0.38 3:(1,2,3,2):0.29	1:(1,3,6,3):0.62 2:(0,3,5,3):0.29 3:(1,3,4,3):0.26	1:(1,4,5,4):0.62 2:(0,4,5,4):0.39 3:(0,4,4,4):0.33
3	(14)11.nen.rpicol.ompcrit.cust.0K30	sing.actual	1:(1,0,0) 1.00 2:(2,0,0) 1.00 3:(5,0,0) 1.00	1:(2,1,3,1):0.00 2:(0,1,4,1):0.00 3:(0,1,5,1):0.00	1:(3,2,6,2):0.71 2:(1,2,3,2):0.68 3:(1,2,6,2):0.47	1:(0,3,2,3):0.67 2:(4,3,7,3):0.68 3:(6,3,7,3):0.46	1:(6,4,7,4):0.68 2:(0,4,4,4):0.40 3:(2,4,6,4):0.40
4	(14)11.nen.rpicol.ompcrit.cust.0K30	sing.log10	1:(1,0,0) 1.00 2:(2,0,0) 1.00 3:(5,0,0) 1.00	1:(2,1,3,1):0.00 2:(0,1,4,1):0.00 3:(0,1,5,1):0.00	1:(5,2,7,2):0.67 2:(1,2,3,2):0.67 3:(1,2,6,2):0.50	1:(0,3,2,3):0.40 2:(1,3,6,3):0.40 3:(0,3,5,3):0.27	1:(5,4,6,4):0.58 2:(0,4,5,4):0.58 3:(1,5,5,4):0.40
5	(14)11.nen.rpicol.ompcrit.cust.0K30	sing.orig	1:(1,0,0) 1.00 2:(2,0,0) 1.00 3:(5,0,0) 1.00	1:(2,1,3,1):0.00 2:(0,1,4,1):0.00 3:(0,1,5,1):0.00	1:(3,2,6,2):0.43 2:(0,2,2,2):0.33 3:(1,2,3,2):0.33	1:(0,3,5,3):0.33 2:(0,3,2,3):0.33 3:(1,3,6,3):0.33	1:(5,4,6,4):0.43 2:(0,4,5,4):0.43 3:(1,4,5,4):0.33
6	-	Average	1.000	0.000	0.441	0.403	0.400
7	-	Mean	1.000	0.000	0.441	0.403	0.400
8	-	Min	1.000	0.000	0.286	0.262	0.333
9	-	Max	1.000	0.000	0.714	0.667	0.714

Figure 12. Recommendation Table Bug 1




Figure 13. diffNLR of process 2 and process 6 buggy(AllReduce() wrong op) vs. bug-free




Figure 14. diffNLR of process 2 and process 6 buggy(AllReduce() wrong size) vs. bug-free




Figure 15. diffNLR of process 1 and process 6 buggy(AllReduce() wrong size) vs. bug-free

42	(7)11.mem.ompcrit.cust.0K10	doub.actual	1:(3_0,4_0):1.00	1:(2_1,3_1):0.00	1:(3_2,6_2):0.62	1:(6_3,7_3):0.88	1:(0_4,4_4):1.00
			2:(2_0,3_0):1.00	2:(0_1,4_1):0.00	2:(3_2,5_2):0.62	2:(4_3,7_3):0.70	2:(4_4,6_4):0.83
			3:(1_0,6_0):1.00	3:(0_1,5_1):0.00	3:(5_2,6_2):0.34	3:(1_3,4_3):0.57	3:(4_4,7_4):0.57
43	(7)11.mem.ompcrit.cust.0K10	doub.log10	1:(3_0,4_0):1.00	1:(2_1,3_1):0.00	1:(3_2,4_2):0.83	1:(1_3,4_3):0.83	1:(0_4,4_4):1.00
			2:(2_0,3_0):1.00	2:(0_1,4_1):0.00	2:(4_2,7_2):0.71	2:(0_3,4_3):0.83	2:(4_4,6_4):0.83
			3:(1_0,6_0):1.00	3:(0_1,5_1):0.00	3:(4_2,4_2):0.71	3:(2_3,4_3):0.71	3:(4_4,5_4):0.83
44	(7)11.mem.ompcrit.cust.0K10	doub.orig	1:(3_0,4_0):1.00	1:(2_1,3_1):0.00	1:(3_2,4_2):0.83	1:(1_3,4_3):0.83	1:(0_4,4_4):1.00
			2:(2_0,3_0):1.00	2:(0_1,4_1):0.00	2:(4_2,7_2):0.71	2:(0_3,4_3):0.83	2:(4_4,6_4):0.83
			3:(1_0,6_0):1.00	3:(0_1,5_1):0.00	3:(1_2,4_2):0.71	3:(2_3,4_3):0.71	3:(4_4,5_4):0.83
45	(7)11.mem.ompcrit.cust.0K10	sing.actual	1:(3_0,4_0):0.75	1:(2_1,3_1):0.00	1:(1_2,4_2):0.66	1:(1_3,7_3):0.89	1:(4_4,6_4):0.52
			2:(2_0,3_0):0.75	2:(0_1,4_1):0.00	2:(4_2,6_2):0.66	2:(6_3,7_3):0.89	2:(2_4,6_4):0.50
			3:(1_0,6_0):0.75	3:(0_1,5_1):0.00	3:(0_1,5_1):0.00	3:(4_3,6_3):0.50	3:(1_4,2_4):0.50
46	(7)11.mem.ompcrit.cust.0K10	sing.log10	1:(3_0,4_0):0.75	1:(2_1,3_1):0.00	1:(4_2,7_2):0.63	1:(0_3,4_3):0.63	1:(4_4,6_4):0.67
			2:(2_0,3_0):0.75	2:(0_1,4_1):0.00	2:(0_2,4_2):0.63	2:(4_3,7_3):0.63	2:(2_4,4_4):0.63
			3:(1_0,6_0):0.75	3:(0_1,5_1):0.00	3:(4_2,5_2):0.63	3:(4_3,5_3):0.63	3:(1_4,4_4):0.50
47	(7)11.mem.ompcrit.cust.0K10	sing.orig	1:(3_0,4_0):0.75	1:(2_1,3_1):0.00	1:(4_2,7_2):0.63	1:(0_3,4_3):0.63	1:(4_4,6_4):0.67
			2:(2_0,3_0):0.75	2:(0_1,4_1):0.00	2:(0_2,4_2):0.63	2:(4_3,7_3):0.63	2:(2_4,4_4):0.63
			3:(1_0,6_0):0.75	3:(0_1,5_1):0.00	3:(4_2,5_2):0.63	3:(4_3,5_3):0.63	3:(1_4,4_4):0.50

Figure 16. Part of ranking table for MisCrit 1-all

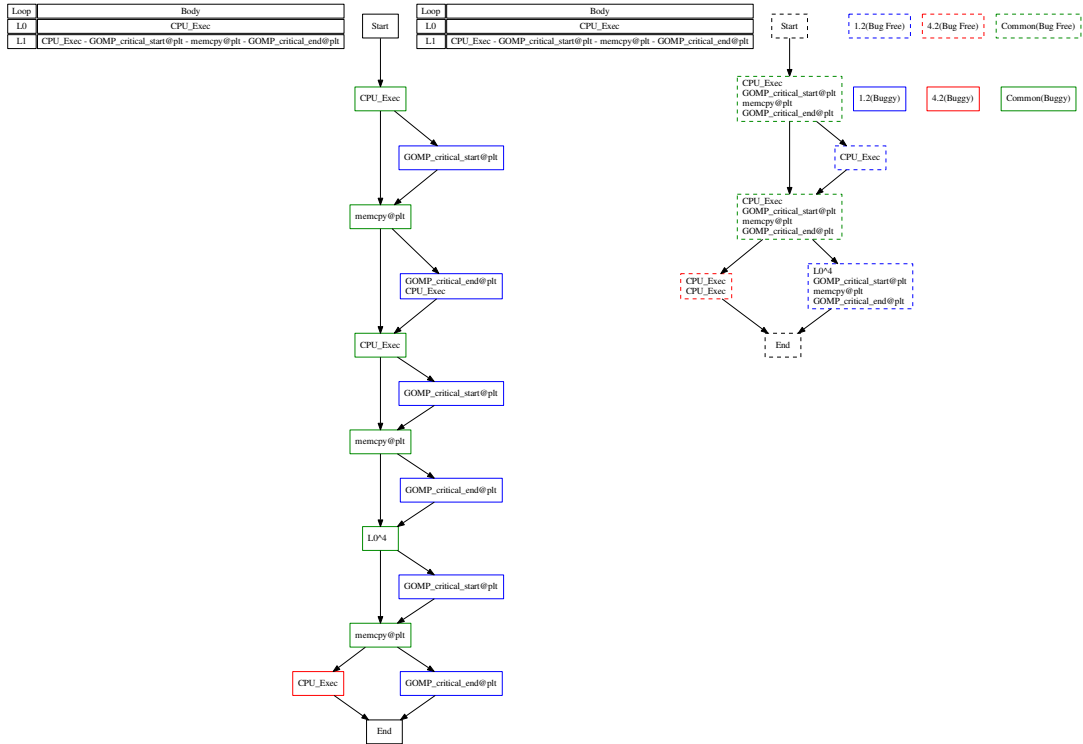


Figure 17. diffNLR of process 1 thread 3 and process 2 thread 3 buggy(missing critical section) vs. bug-free

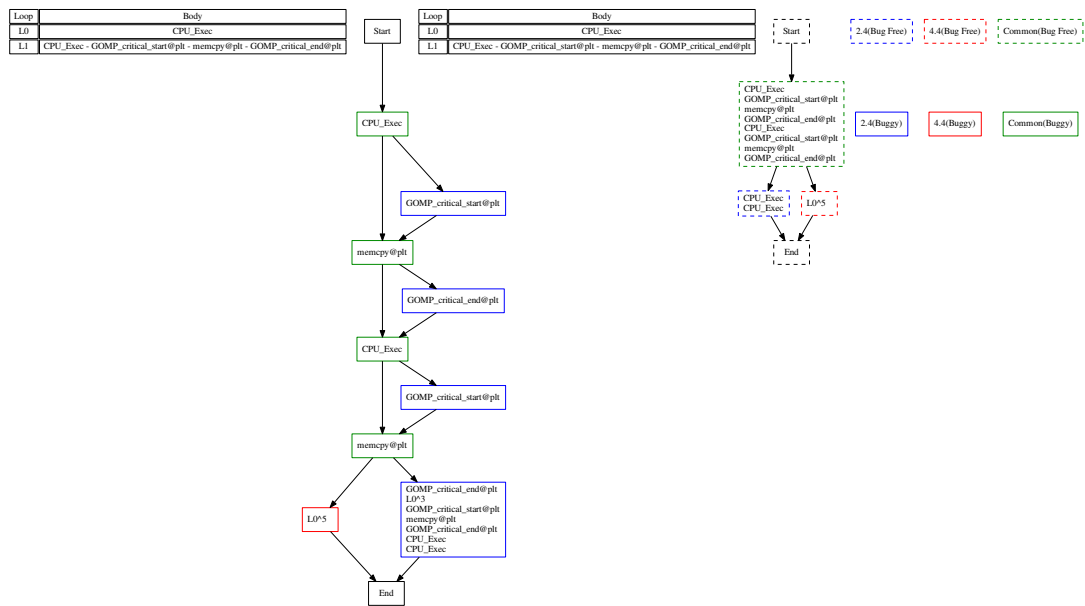


Figure 18. diffNLR of process 1 thread 3 and process 2 thread 3 buggy(missing critical section) vs. bug-free